

Machine Learning

Summer Semester 2019, Homework 2

Prof. Dr. J. Peters, H. Abdulsamad, S. Stark, D. Koert



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Total points: 90 + 15 bonus

Due date: Friday, 14 June 2019 17:00

Hand in a PDF over Moodle and a printed version to the postbox at (S2/02 | E315)

Name, Surname, ID Number

Steffen Schäfer, 2635897

Peter Nickl, 1941346

Problem 2.1 Bayesian Decision Theory [20 Points]

In this exercise, we consider data generated by a mixture of two Gaussian distributions with parameters $\{\mu_1, \sigma_1\}$ and $\{\mu_2, \sigma_2\}$. Each Gaussian represents a class labeled C_1 and C_2 , respectively.

a) Optimal Boundary [4 Points]

Explain in one short sentence what Bayesian Decision Theory is. What is its goal? What condition does hold at the optimal decision boundary? When do we decide for class C_1 over C_2 ?

b) Decision Boundaries [8 Points]

If both classes have equal prior probabilities $p(C_1) = p(C_2)$ and the same variance $\sigma_1 = \sigma_2$, derive the decision boundary x^* analytically as a function of the two means μ_1 and μ_2 .

c) Different Misclassification Costs [8 Points]

Assume $\mu_1 > 0$, $\mu_1 = 2\mu_2$, $\sigma_1 = \sigma_2$ and $p(C_1) = p(C_2)$. If misclassifying sample $x \in C_2$ as class C_1 is three times more expensive than the opposite, how does the decision boundary change? Derive the boundary analytically. (There is no cost for correctly classifying samples.)

Problem 2.2 Density Estimation [30 Points + 15 Bonus]

In this exercise, you will use the datasets `densEst1.txt` and `densEst2.txt`. The datasets contain 2D data belonging to two classes, C_1 and C_2 .

a) Gaussian Maximum Likelihood Estimation [10 Points]

Derive the ML estimate for the mean and covariance of the **multivariate** Gaussian distribution. Start your derivations with the function you optimize. Assume that you can collect i.i.d data. (Hint: you can find many matrix identities on the Matrix Cookbook (<https://www.math.uwaterloo.ca/~hwolkowi/matrixcookbook.pdf>) and at http://en.wikipedia.org/wiki/Matrix_calculus.)

The Multivariate Gaussian Distribution is defined as

$$f_x(\mu|\Sigma) = \frac{1}{\sqrt{(2\pi)^p \det(\Sigma)}} \exp\left(-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)\right) \quad (23)$$

Σ is the Covariance Matrix defined as:

$$\begin{pmatrix} \sigma_{xx}^2 & \sigma_{xy}^2 \\ \sigma_{xy}^2 & \sigma_{yy}^2 \end{pmatrix} \quad (24)$$

The resulting likelihood function of $f_x(x)$ is:

$$L(\mu, \Sigma) = \prod_{n=1}^N f_x(\mu|\Sigma) \quad (25)$$

We now take the log of the likelihood function.

(Doesnt matter if we use theta or mu and sigma)

$$\log L(\mu, \Sigma) = \sum_{n=1}^N \log f_x(\mu|\Sigma) \quad (26)$$

Resulting in the log likelihood function for a multivariate gaussian distribution.

$$\log L(\mu, \Sigma) = \sum_{n=1}^N \log\left(\frac{1}{\sqrt{(2\pi)^p \det(\Sigma)}} \exp\left(-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)\right)\right) \quad (27)$$

Now we want to derive the mean and covariance for this distribution. Starting with the mean, we need to maximize the partial equation.

$$\frac{\partial L}{\partial \mu} = \sum_{n=1}^N \underbrace{\log\left(\frac{1}{\sqrt{(2\pi)^p \det(\Sigma)}}\right)'}_0 + \sum_{n=1}^N \left(-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)\right)' \quad (28)$$

For matrix Calculus we use following rules:

$$\frac{\partial (Xa+b)^T C (Xa+b)}{\partial x} = (C + C^T)(Xa+b)a^T \quad (29)$$

With $X = x$, $a = 1$, thus $a^T = 1$ and $b = -\mu$. We get

$$\sum_{n=1}^N \left(\Sigma + (\Sigma^{-1})^T (x - \mu) \right) = 0 \quad (30)$$

Since Σ is a symmetrical Matrix $\Sigma + \Sigma^T = 2 \cdot \Sigma$ holds true. This matrix is also constant, meaning its independent on the sum and can be written before it.

$$2\Sigma^{-1} \cdot \sum_{n=1}^N (x - \mu) = 0 \quad (31)$$

$$\sum_{n=1}^N x_i = \sum_{n=1}^N \mu \quad (32)$$

Since μ is a constant we can write the sum over a constant as $N \cdot \mu$

$$\sum_{n=1}^N x_i = N \cdot \mu \quad (33)$$

Resulting in the following equation for the mean:

$$\mu = \frac{1}{N} \sum_{n=1}^N x_i \quad (34)$$

Now we want to derive the maximum Likelihood estimate for the covariance of the multivariate Gaussian distribution. Again we Start with Equation (27). For this derivation we will use the following rules for matrix calculus.

$$\frac{\partial}{\partial A} \log(\det(A)) = (A^{-1})^T \quad (35)$$

$$\frac{\partial}{\partial A} x^T A x = x x^T \quad (36)$$

With $A = \Sigma$ and $x = (x - \mu)$. First we split the log of gaussian equation up into three part for easier derivation.

$$\frac{\partial L}{\partial \Sigma} = \underbrace{\log\left(\frac{1}{\sqrt{2\pi}^p}\right)}_a + \sum_{n=1}^N \left(\underbrace{\log\left(\frac{1}{\sqrt{\det(\Sigma)}}\right)}_b + \underbrace{\log \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1} (x - \mu)\right)}_c \right) \quad (37)$$

For a:

$$\frac{\partial L_a}{\partial \Sigma} = 0 \quad (38)$$

For b:

$$\log\left(\frac{1}{\sqrt{\det(\Sigma)}}\right) = \log\left(\frac{1}{\det(\Sigma)^{\frac{1}{2}}}\right) = \frac{1}{2} \cdot \log\left(\frac{1}{\det(\Sigma)}\right) \quad (39)$$

$$\frac{\partial L_b}{\partial \Sigma} = \frac{1}{2} \cdot \frac{1}{(\Sigma^{-1})^T} = \frac{1}{2} \cdot \Sigma \quad (40)$$

For c:

$$\frac{\partial L_c}{\partial \Sigma} = -\frac{1}{2} (x - \mu) \cdot (x - \mu)^T \quad (41)$$

In total we get the following derivative:

$$\frac{\partial L}{\partial \Sigma} = \sum_{n=1}^N \frac{1}{2} \Sigma - \sum_{n=1}^N \left(\frac{1}{2} (x - \mu) \cdot (x - \mu)^T \right) = 0 \quad (42)$$

Multiplied by $\frac{1}{2}$:

$$\sum_{n=1}^N \Sigma = \sum_{n=1}^N ((x - \mu) \cdot (x - \mu)^T) \quad (43)$$

Since Σ is constant we can write $\sum_{n=1}^N \Sigma = N \cdot \Sigma$. Resulting in the final equation for the covariance

$$\Sigma = \frac{1}{N} \sum_{n=1}^N (x_i - \mu) \cdot (x_i - \mu)^T \quad (44)$$

b) Prior Probabilities [2 Points]

Compute the prior probability of each class from the dataset.

c) Biased ML Estimate [5 Points]

Define the bias of an estimator and write how we can compute it. Then calculate the biased and unbiased estimates of the conditional distribution $p(x|C_i)$, assuming that each class can be modeled with a Gaussian distribution. Which parameters have to be calculated? Show the final result and attach a snippet of your code. Do not use existing functions, but rather implement the computations by yourself!

d) Class Density [5 Points]

Using the unbiased estimates from the previous question, fit a Gaussian distribution to the data of each class. Generate a single plot showing the data points and the probability densities of each class. (Hint: use the contour function for plotting the Gaussians.)

e) Posterior [8 Points]

In a single graph, plot the posterior distribution of each class $p(C_i|x)$ and show the decision boundary.

f) Bayesian Estimation [15 Bonus Points]

State the generic case of Bayesian linear regression with data $\langle X, Y \rangle$ and parameters θ . What do we assume about the data, the model and the parameters?

Formulate the posterior distribution for your model parameters given the data, i.e., $p(\theta|X, Y)$, and derive its mean and covariance, assuming that the model of the output variable is a Gaussian distribution with a fixed variance.

What do we do when we want to predict a new point?

Which are the advantages of being Bayesian?

Problem 2.3 Non-parametric Density Estimation [20 Points]

In this exercise, you will use the datasets `nonParamTrain.txt` for training and `nonParamTest.txt` for evaluating the performance of your model.

a) Histogram [4 Points]

Compute and plot the histograms using 0.02, 0.5, 2.0 size bins of the training data. Intuitively, indicate which bin size performs the best and explain why. Knowing only these three trials, would you be sure that the one you picked is truly the best one? Attach the plot of the histograms and a snippet of your code.

b) Kernel Density Estimate [6 Points]

Compute the probability density estimate using a Gaussian kernel with $\sigma = 0.03$, $\sigma = 0.2$ and $\sigma = 0.8$ of the training data. Compute the log-likelihood of the data for each case, compare them and show which parameter performs the best. Generate a single plot with the different density estimates and attach it to your solutions. Plot the densities in the interval $x \in [-4, 8]$, attach a snippet of your code and comment the results.

c) K-Nearest Neighbors [6 Points]

Estimate the probability density with the K-nearest neighbors method with $K = 2, K = 8, K = 35$. Generate a single plot with the different density estimates and attach it to your solutions. Plot the densities in the interval $x \in [-4, 8]$, attach a snippet of your code and comment the results.

d) Comparison of the Non-Parametric Methods [4 Points]

Estimate the log-likelihood of the testing data using the KDE estimators and the K-NN estimators. Why do we need to test them on a different data set? Compare the log-likelihoods of the estimators w.r.t. both the training and testing sets in a table. Which estimator would you choose?

Problem 2.4 Expectation Maximization [20 Points]

In this exercise, you will use the datasets `gmm.txt`. It contains data from a Gaussian Mixture Model with four 2-dimensional Gaussian distributions.

a) Gaussian Mixture Update Rules [2 Points]

Define the model parameters and the update rules for your model. Specify the E- and M-steps of the algorithm.

b) EM [18 Points]

Implement the Expectation Maximization algorithm for Gaussian Mixture Models. Initialize your model uniformly. Generate plots at different iterations $t_i \in [1, 3, 5, 10, 30]$, showing the data and the mixture components, and plot the log-likelihood for every iteration $t_i = 1 : 30$. Attach a snippet of your code.